

Budgeted Reinforcement Learning in Continuous State Space

Nicolas Carrara ³ Edouard Leurent^{3,4} Tanguy Urvoy ¹
Romain Laroche ² Odalric-Ambrym Maillard ³ Olivier
Pietquin ^{3,4}

¹Orange Labs

²Microsoft Montréal.

³Univ. Lille, CNRS, Centrale Lille, INRIA UMR 9189 - CRISTAL

⁴Google Research, Brain Team, Paris

⁴Renault

10 august 2018.


Introduction — Use-case




Introduction — Use-case




Introduction — Use-case




Introduction — Use-case



Introduction — Use-case



Introduction — Use-case



Introduction — Use-case

Problem

Given past trajectories, find a way to:

- ▶ gather gold as much as possible;
- ▶ limit the villager casualties under some budget;
- ▶ being able to change the budget in real time.

Solution

- ▶ This problem can be cast as a Budgeted Markov Decision Process.

Setting

Markov Decision Process

We define a MDP as a tuple $(\mathcal{S}, \mathcal{A}, P, R_r, \gamma)$ where:

- ▶ \mathcal{S} is the state space, \mathcal{A} the action space,
- ▶ $R_r \in \mathbb{R}^{\mathcal{S} \times \mathcal{A}}$ the rewards,
- ▶ $P \in \mathcal{M}(\mathcal{S})^{\mathcal{S} \times \mathcal{A}}$ the dynamics,
- ▶ and γ the discounted factor.

Objective

- ▶ $G_r^\pi = \sum_{t=0}^{\infty} \gamma^t R_r(s_t, a_t)$ the γ -discounted return of rewards.
- ▶ Find π^* s.t $\forall s \in \mathcal{S}$:

$$\pi^* \in \arg \max_{\pi \in \mathcal{M}(\mathcal{A})^{\mathcal{S}}} \mathbb{E}[G_r^\pi | s_0 = s] \quad (1)$$

Setting

Constrained Markov Decision Process

We define a **CMDP** as a tuple $(\mathcal{S}, \mathcal{A}, P, R_r, R_c, \gamma, \beta)$ where:

- ▶ \mathcal{S} is the state space, \mathcal{A} the action space,
- ▶ $R_r \in \mathbb{R}^{\mathcal{S} \times \mathcal{A}}$ the rewards, and $R_c \in \mathbb{R}^{\mathcal{S} \times \mathcal{A}}$ the costs
- ▶ $P \in \mathcal{M}(\mathcal{S})^{\mathcal{S} \times \mathcal{A}}$ the dynamics,
- ▶ γ the discounted factor, and β the budget.

Objective

- ▶ $G_r^\pi = \sum_{t=0}^{\infty} \gamma^t R_r(s_t, a_t)$ the γ -discounted return of rewards.
- ▶ $G_c^\pi = \sum_{t=0}^{\infty} \gamma^t R_c(s_t, a_t)$ the γ -discounted return of costs.
- ▶ Find π^* s.t. $\forall s \in \mathcal{S}$:

$$\begin{aligned} \pi^* &\in \arg \max_{\pi \in \mathcal{M}(\mathcal{A})^{\mathcal{S}}} \mathbb{E}[G_r^\pi | s_0 = s] \\ \text{s.t. } &\mathbb{E}[G_c^\pi | s_0 = s] \leq \beta \end{aligned} \tag{2}$$

Setting

Budgeted Markov Decision Process

We define a **BMDP** as a tuple $(\mathcal{S}, \mathcal{A}, P, R_r, R_c, \gamma, \mathcal{B})$ where:


- ▶ \mathcal{S} is the state space, \mathcal{A} the action space,
- ▶ $R_r \in \mathbb{R}^{\mathcal{S} \times \mathcal{A}}$ the rewards, and $R_c \in \mathbb{R}^{\mathcal{S} \times \mathcal{A}}$ the costs
- ▶ $P \in \mathcal{M}(\mathcal{S})^{\mathcal{S} \times \mathcal{A}}$ the dynamics,
- ▶ γ the discounted factor, and \mathcal{B} the budget space.

Objective


- ▶ $G_r^\pi = \sum_{t=0}^{\infty} \gamma^t R_r(s_t, a_t)$ the γ -discounted return of rewards.
- ▶ $G_c^\pi = \sum_{t=0}^{\infty} \gamma^t R_c(s_t, a_t)$ the γ -discounted return of costs.
- ▶ Find π^* s.t. $\forall (s, \beta) \in \mathcal{S} \times \mathcal{B}$:

$$\begin{aligned} \pi^* \in & \arg \max_{\pi \in \mathcal{M}(\mathcal{A} \times \mathcal{B})^{\mathcal{S} \times \mathcal{B}}} \mathbb{E}[G_r^\pi | s_0 = s, \beta_0 = \beta] \\ \text{s.t. } & \mathbb{E}[G_c^\pi | s_0 = s, \beta_0 = \beta] \leq \beta \end{aligned} \tag{3}$$

The Dynamic Programming solution: intuition




The Dynamic Programming solution: intuition



$$r = 10$$
$$c = 2$$

The Dynamic Programming solution: intuition



$$r = 2 + 9\beta_1/2$$

$$c = 1 + \beta_1/2$$

$$r = 10$$

$$c = 2$$

The Dynamic Programming solution: intuition

$$\rho_0 = \arg \max_{\rho, \beta_1} \rho 10 + (1 - \rho)(2 + 9\beta_1)$$


$$\text{s.t } \rho 2 + (1 - \rho)(1 + \beta_1/2) \leq \beta$$

$$r = 2 + 9\beta_1/2$$

$$c = 1 + \beta_1/2$$

$$r = 10$$

$$c = 2$$



The Dynamic Programming solution: intuition

$$\rho_0 = \arg \max_{\rho, \beta_1} \rho 10 + (1 - \rho)(2 + 9\beta_1)$$

$$\text{s.t. } \rho 2 + (1 - \rho)(1 + \beta_1/2) \leq \beta$$

$$r = \rho_0 10 + (1 - \rho_0)(2 + 9\beta_1)$$


$$c = \rho_0 2 + (1 - \rho_0)(1 + \beta_1/2)$$

$$r = 2 + 9\beta_1/2$$

$$c = 1 + \beta_1/2$$

$$r = 10$$

$$c = 2$$



Augmented Settings

Budgeted policies π

- ▶ Take a budget β as an additional input
- ▶ Output a next budget β'
- ▶ $\pi : \underbrace{(s, \beta)}_{\bar{s}} \rightarrow \underbrace{(a, \beta')}_{{\bar{a}}}$

Domain

- ▶ States $\bar{\mathcal{S}} = \mathcal{S} \times \mathcal{B}$.
- ▶ Actions $\bar{\mathcal{A}} = \mathcal{A} \times \mathcal{B}$.
- ▶ Dynamics $\bar{P}((s', \beta') \mid (s, \beta), (a, \beta_a)) \stackrel{\text{def}}{=} P(s' \mid s, a) \delta(\beta' - \beta_a)$.

2D signals

1. Rewards $R = (R_r, R_c)$
2. Returns $G^\pi = (G_r^\pi, G_c^\pi)$
3. $V^\pi(\bar{s}) = (V_r^\pi, V_c^\pi) \stackrel{\text{def}}{=} \mathbb{E}[G^\pi \mid \bar{s}_0 = \bar{s}]$
4. $Q^\pi(\bar{s}, \bar{a}) = (Q_r^\pi, Q_c^\pi) \stackrel{\text{def}}{=} \mathbb{E}[G^\pi \mid \bar{s}_0 = \bar{s}, \bar{a}_0 = \bar{a}]$

Policy Evaluation

The Bellman Expectation equations are preserved, and the Bellman Expectation Operator \mathcal{T}^π is a γ -contraction.

Augmented Optimality

Definition

In that order, we want to:

- (i) Respect the budget β :

$$\Pi_a(\bar{s}) \stackrel{\text{def}}{=} \{\pi \in \Pi : V_c^\pi(s, \beta) \leq \beta\}$$

- (ii) Maximise the rewards:

$$V_r^*(\bar{s}) \stackrel{\text{def}}{=} \max_{\pi \in \Pi_a(\bar{s})} V_r^\pi(\bar{s})$$

$$\Pi_r(\bar{s}) \stackrel{\text{def}}{=} \arg \max_{\pi \in \Pi_a(\bar{s})} V_r^\pi(\bar{s})$$

- (iii) Minimise the costs:

$$V_c^*(\bar{s}) \stackrel{\text{def}}{=} \min_{\pi \in \Pi_r(\bar{s})} V_c^\pi(\bar{s}),$$

$$\Pi^*(\bar{s}) \stackrel{\text{def}}{=} \arg \min_{\pi \in \Pi_r(\bar{s})} V_c^\pi(\bar{s})$$

We define the budgeted action-value function Q^* similarly

Budgeted Bellman Optimality Equation

Theorem (Budgeted Bellman Optimality Equation)

Q^* verifies the following equation:

$$Q^*(\bar{s}, \bar{a}) = \mathcal{T} Q^*(\bar{s}, \bar{a}) \\ \stackrel{\text{def}}{=} R(\bar{s}, \bar{a}) + \gamma \sum_{\bar{s}' \in \bar{\mathcal{S}}} \bar{P}(\bar{s}' | \bar{s}, \bar{a}) \sum_{\bar{a}' \in \bar{\mathcal{A}}} \pi_{\text{greedy}}(\bar{a}' | \bar{s}'; Q^*) Q^*(\bar{s}', \bar{a}'),$$

where the greedy policy π_{greedy} is defined by:

$$\pi_{\text{greedy}}(\bar{a} | \bar{s}; Q) \in \arg \min_{\rho \in \Pi_r^Q} \mathbb{E}_{\bar{a} \sim \rho} Q_c(\bar{s}, \bar{a}),$$

$$\text{where } \Pi_r^Q \stackrel{\text{def}}{=} \arg \max_{\rho \in \mathcal{M}(\bar{\mathcal{A}})} \mathbb{E}_{\bar{a} \sim \rho} Q_r(\bar{s}, \bar{a})$$

$$\text{s.t. } \mathbb{E}_{\bar{a} \sim \rho} Q_c(\bar{s}, \bar{a}) \leq \beta$$

Optimality of the policy

Proposition (Optimality of the policy)

$\pi_{greedy}(\cdot ; Q^*)$ is *simultaneously optimal* in all states $\bar{s} \in \bar{\mathcal{S}}$:

$$\pi_{greedy}(\cdot ; Q^*) \in \Pi^*(\bar{s})$$


In particular, $V^{\pi_{greedy}(\cdot ; Q^)} = V^*$ and $Q^{\pi_{greedy}(\cdot ; Q^*)} = Q^*$.*

Solving the untractable program


Proposition ($\pi_{\text{greedy}} = \pi_{\text{hull}}$)

π_{greedy} can be computed explicitly, as a mixture of two points that lie on the convex hull of Q .


Solving the non-linear programming problem: intuition




Solving the non-linear programming problem: intuition




Solving the non-linear programming problem: intuition



Solving the non-linear programming problem: intuition



Solving the non-linear programming problem: intuition




Not a contraction

Theorem (Non-Contractivity)


For any BMDP $(\mathcal{S}, \mathcal{A}, P, R_r, R_c, \gamma)$ with $|\mathcal{A}| \geq 2$, \mathcal{T} is not a contraction.

$$\forall \varepsilon > 0, \exists Q^1, Q^2 \in (\mathbb{R}^2)^{\overline{\mathcal{SA}}} : \|\mathcal{T}Q^1 - \mathcal{T}Q^2\|_{\infty} \geq \frac{1}{\varepsilon} \|Q^1 - Q^2\|_{\infty}$$


Not a contraction: intuition



Not a contraction: intuition



Not a contraction: intuition



Contractivity on smooth Q -functions

Conjecture (Contractivity \mathcal{L}_γ)

\mathcal{T} is a contraction when restricted to the subset \mathcal{L}_γ of Q -functions such that " Q_r is L -Lipschitz with respect to Q_c ", with $L < \frac{1}{\gamma} - 1$.

Budgeted Dynamic Programming

Algorithm 1: Budgeted Value-Iteration

Data: P, R_r, R_c

Result: Q^*

- 1 $Q_0 \leftarrow 0$
 - 2 **repeat**
 - 3 $Q_{k+1} \leftarrow \mathcal{T}Q_k$
 - 4 **until** *convergence*
-

Budgeted Reinforcement Learning

We address several limitations of Budgeted Value-Iteration

1. If the P , R_r and R_c are unknown:

- ▶ Work with a batch of samples $\mathcal{D} = \{(\bar{s}_i, \bar{a}_i, r_i, \bar{s}'_i)\}_{i \in [0, N]}$
- ▶ Replace \mathcal{T} with a sampling operator $\hat{\mathcal{T}}$:

$$\hat{\mathcal{T}}Q(\bar{s}_i, \bar{a}_i, r_i, \bar{s}'_i) \stackrel{\text{def}}{=} r_i + \gamma \sum_{\bar{a}'_i \in \mathcal{A}_i} \pi_{\text{greedy}}(\bar{a}'_i | \bar{s}'_i; Q) Q(\bar{s}'_i, \bar{a}'_i).$$

2. If \mathcal{S} is continuous:

- ▶ Employ function approximation Q_θ , and minimise a regression loss

$$\mathcal{L}(Q_\theta, Q_{\text{target}}; \mathcal{D}) = \sum_{\mathcal{D}} \|Q_\theta(\bar{s}, \bar{a}) - Q_{\text{target}}(\bar{s}, \bar{a}, r, \bar{s}')\|_2^2$$

Budgeted Fitted-Q

Algorithm 2: Budgeted Fitted-Q Iteration


Data: \mathcal{D}

Result: Q^*

- 1 $Q_{\theta_0} \leftarrow 0$
 - 2 **repeat**
 - 3 $\theta_{k+1} \leftarrow \arg \min_{\theta} \mathcal{L}(Q_{\theta}, \hat{\mathcal{T}} Q_{\theta_k}; \mathcal{D})$
 - 4 **until** *convergence*
-

More scaling

- ▶ CPU parallel computing of the targets
$$\sum_{\overline{a'_i} \in \mathcal{A}_i} \pi_{\text{greedy}}(\overline{a'_i} | \overline{s'_i}; Q) Q(\overline{s'_i}, \overline{a'_i}) \quad \forall i$$
- ▶ Same for samples generation.
- ▶ Neural Network as function approximator:




Experiments: performances of BFTQ

- ▶ Baseline: λ -FTQ, Lagrangian relaxation
 - ▶ $R_r(s, a) \leftarrow R_r(s, a) - \lambda R_c(s, a)$ where $\lambda \geq 0$
- .
- ▶ Applications:
 - ▶ dialogue systems
 - ▶ autonomous driving

Experiments: dialogue systems

- ▶ A slot-filling problem: the agent (the dialogue system) fills a form by asking the user each slot.
- ▶ Two ways to deal with recognition errors:
 - ▶ ask to repeat with voice (safe/slow),
 - ▶ Ask to repeat with numeric pad (unsafe/fast).


Experiments: dialogue systems



Experiments: autonomous driving

- ▶ the agent (the car) is on a two-way road with a car in front of it:
 - ▶ it can stay behind (safe/slow),
 - ▶ it can overtake (unsafe/fast).

Experiments: autonomous driving



Experiments: autonomous driving

► BFTQ on the highway environment

Risk-sensitive exploration

How to collect the batch \mathcal{D} ?

- ▶ We propose an ε -greedy exploration procedure
 - ▶ Sample an initial budget β_0
 - ▶ At each step, where $\bar{s} = (s, \beta)$ only explore feasible budgets:

$$\bar{a} = (a, \beta_a) \sim \mathcal{U}(\Delta_{\mathcal{AB}})$$


where Δ is s.t. $\mathbb{P}(a, \beta_a | s, \beta)$ verifies $\mathbb{E}[\beta_a] \leq \beta$

Experiments: risk-sensitive exploration



- ▶ Validate the risk-sensitive exploration procedure on the corridor environment
- ▶ Learn 2 BFTQ policies with respectively:
 - ▶ A batch generated by a risk-neutral ε -greedy procedure
 - ▶ A batch generated by a risk-sensitive ε -greedy procedure

Experiments: corridors

- ▶ 2 corridors:
 - ▶ high costs/high rewards around the starting state
 - ▶ no costs/low rewards around the starting state
- ▶ The outermost cell is the one yielding the most reward



Experiments: corridors



Experiments: corridors

► Risk-sensitive vs Risk-Neutral on the corridors environment

Summary

- + Budgeted Bellman Optimality Operator.
 - ▶ Fixed point.
 - ▶ Not a contraction but converging in practice.
- + Scalable for RL in continuous state space.
 - ▶ Function approximation with Neural Network (dedicated architecture).
 - ▶ Solving of the untractable program using convex hull.
 - ▶ CPU parallel computing of the target.
 - ▶ Risk-sensitive exploration procedure.
- + Experiments on two applications.
 - ▶ BFTQ reaches similar performances as Lagrangian relaxation,
 - ▶ with no need for calibration,
 - ▶ and less variance.